

## ▼ Working Jupyter Notebook for HW2:

<https://colab.research.google.com/drive/1CqsTUeYjTt-WCwewNb7eVhNnV5Jaou-I?usp=sharing>

```
# %load_ext tensorboard
%reload_ext tensorboard
```

## ▼ Problem 1

### ▼ Tensorboard Setup

```
from datetime import datetime
from pathlib import Path

from torch.utils.tensorboard import SummaryWriter

writer = SummaryWriter("runs/LeNet5_experiment")

datetime_str = datetime.now().strftime('%b%d_%H-%M-%S')

base_folder = Path('runs/LeNet5_experiment')
base_folder.mkdir(parents=True, exist_ok=True)

logging_dir = base_folder / Path(datetime_str)
logging_dir.mkdir(exist_ok=True)

logging_dir = str(logging_dir.absolute())

writer = SummaryWriter(log_dir=logging_dir)

%tensorboard --logdir {logging_dir} --port 6000
```

Filter runs (regex)

Filter tags (regex)

All   Scalars   Image   Histogram

☒ Run

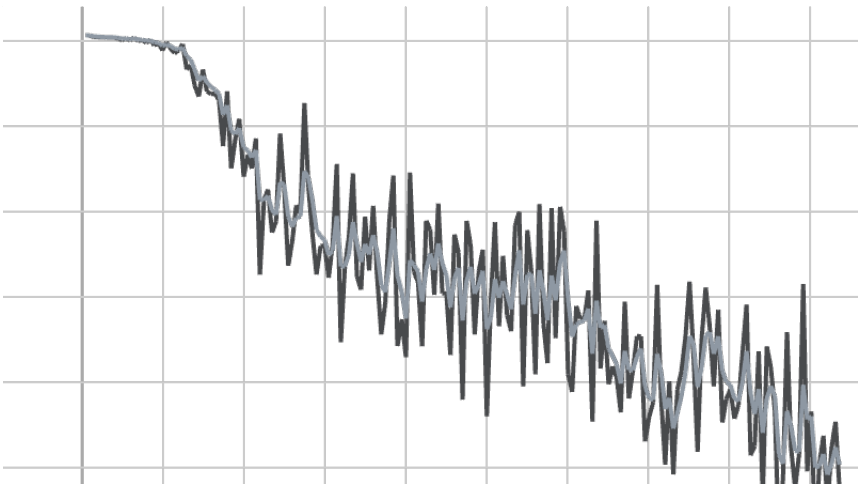
☒ .

Pinned

Pin cards for a quick view and comparison

Loss

Loss/train\_batch



histogram 28 cards

Previous

Next

histogram/test\_conv1\_biases

▼ Importing Libraries

```
import matplotlib.pyplot as plt
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
```

## ▼ Define Variables

```
# Initialize dictionary to save activations
# activations = {}

logging_interval = 10
batch_size = 128
epochs = 10

train_accuracies = []
test_accuracies = []
train_losses = []
```

## ▼ Initializing Functions

```
def rgb_to_grayscale(img):
    gray = 0.299 * img[0] + 0.587 * img[1] + 0.114 * img[2]
    return gray.unsqueeze(0) # Adds a channel at 0th dimension

activations = {}
def hook_fn(module, input, output):
    # layer_name = str(module)
    activations[str(module)] = output
```

## ▼ Define LeNet5 Architecture

```

class LeNet5(nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 5)
        self.conv2 = nn.Conv2d(32, 64, 5)
        self.fc1 = nn.Linear(5 * 5 * 64, 1024)
        self.fc2 = nn.Linear(1024, 10)
        self.stats = {}

    def forward(self, x):
        x = nn.functional.relu(self.conv1(x))
        self.stats['conv1_weights'] = self.conv1.weight.data.cpu().numpy()
        self.stats['conv1_biases'] = self.conv1.bias.data.cpu().numpy()
        x = nn.functional.max_pool2d(x, 2)
        x = nn.functional.relu(self.conv2(x))
        self.stats['conv2_weights'] = self.conv2.weight.data.cpu().numpy()
        self.stats['conv2_biases'] = self.conv2.bias.data.cpu().numpy()
        x = nn.functional.max_pool2d(x, 2)
        x = x.view(-1, 5 * 5 * 64)
        x = nn.functional.relu(self.fc1(x))
        self.stats['fc1_weights'] = self.fc1.weight.data.cpu().numpy()
        self.stats['fc1_biases'] = self.fc1.bias.data.cpu().numpy()
        x = self.fc2(x)
        self.stats['fc2_weights'] = self.fc2.weight.data.cpu().numpy()
        self.stats['fc2_biases'] = self.fc2.bias.data.cpu().numpy()
        return nn.functional.softmax(x, dim=1)

model = LeNet5()
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

```

## ▼ Data Loading

```

# Define the transformation for training and test datasets
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Lambda(rgb_to_grayscale),
    transforms.Normalize((0.5,), (0.5,))
])

```

```
# Load the CIFAR10 dataset
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, t
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, t

    Files already downloaded and verified
    Files already downloaded and verified

# DataLoader
train_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle
test_loader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=F
```

## ▼ Training Loop

```
train_losses = []
test accuracies = []
train accuracies = []

for epoch in range(epochs):
    model.train()
    correct = 0
    total = 0
    running_loss = 0.0

    for i, data in enumerate(train_loader, 0):
        inputs, labels = data
        optimizer.zero_grad()

        # Forward pass
        outputs = model(inputs)
        loss = criterion(outputs, labels)

        # Backward pass and optimization
        loss.backward()
        optimizer.step()

        # Calculate accuracy
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        # Calculate running loss
        running_loss += loss.item()
        n_iter = epoch * len(train_loader) + i

        if (i + 1) % logging interval == 0:
```

```

        writer.add_scalar('Loss/train_batch', loss.item(), n_iter)
        writer.add_scalar('Accuracy/Train', 100 * correct / total, epoch)

    for key, value in model.stats.items():
        value_flat = value.flatten()
        writer.add_scalar(f'statistics/train_{key}_mean', np.mean(value_flat))
        writer.add_scalar(f'statistics/train_{key}_std', np.std(value_flat))
        writer.add_scalar(f'statistics/train_{key}_min', np.min(value_flat))
        writer.add_scalar(f'statistics/train_{key}_max', np.max(value_flat))
        writer.add_histogram(f'histogram/train_{key}', value_flat, n_iter)

    for name, param in model.named_parameters():
        if 'weight' in name:
            writer.add_scalar(f'statistics/train_{name}_min', param.min().item())
            writer.add_scalar(f'statistics/train_{name}_max', param.max().item())
            writer.add_scalar(f'statistics/train_{name}_mean', param.mean().item())
            writer.add_scalar(f'statistics/train_{name}_std', param.std().item())
            writer.add_histogram(f'histogram/test_{name}', param, n_iter)
        elif 'bias' in name:
            writer.add_scalar(f'statistics/train_{name}_min', param.min().item())
            writer.add_scalar(f'statistics/train_{name}_max', param.max().item())
            writer.add_scalar(f'statistics/train_{name}_mean', param.mean().item())
            writer.add_scalar(f'statistics/train_{name}_std', param.std().item())
            writer.add_histogram(f'histogram/train_{name}', param, n_iter)

    train_accuracies.append(100 * correct / total)
    train_losses.append(running_loss / len(train_loader))

    print(f"Epoch {epoch+1}, Loss: {running_loss / len(train_loader)}, Accuracy: {100 * correct / total}")

# Test Loop
model.eval()
correct = 0
total = 0

with torch.no_grad():
    for j, data in enumerate(test_loader):

        images, labels = data
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        m_iter = epoch * len(test_loader) + j
        if (j + 1) % logging_interval == 0:
            for key, value in model.stats.items():
                value_flat = value.flatten()
                writer.add_scalar(f'statistics/test_{key}_mean', np.mean(value_flat))

```

```

        writer.add_scalar(f'statistics/test_{key}_std', np.std(value_
        writer.add_scalar(f'statistics/test_{key}_min', np.min(value_
        writer.add_scalar(f'statistics/test_{key}_max', np.max(value_
        writer.add_histogram(f'histogram/test_{key}', value_flat, m_j

for name, param in model.named_parameters():
    if 'weight' in name:
        writer.add_scalar(f'statistics/test_{name}_min', param.mi
        writer.add_scalar(f'statistics/test_{name}_max', param.ma
        writer.add_scalar(f'statistics/test_{name}_mean', param.n
        writer.add_scalar(f'statistics/test_{name}_std', param.st
        writer.add_histogram(f'histogram/test_{name}', param, m_j
    elif 'bias' in name:
        writer.add_scalar(f'statistics/test_{name}_min', param.mi
        writer.add_scalar(f'statistics/test_{name}_max', param.ma
        writer.add_scalar(f'statistics/test_{name}_mean', param.n
        writer.add_scalar(f'statistics/test_{name}_std', param.st
        writer.add_histogram(f'histogram/test_{name}', param, m_j
test accuracies.append(100 * correct / total)
print(f"Test Accuracy: {100 * correct / total}%")
writer.add_scalar('Accuracy/test', 100 * correct / total, epoch)

```

```

Epoch 1, Loss: 2.2880112874843275, Accuracy: 14.786%
Test Accuracy: 21.01%
Epoch 2, Loss: 2.1938467855038852, Accuracy: 25.78%
Test Accuracy: 28.96%
Epoch 3, Loss: 2.15173975829883, Accuracy: 30.342%
Test Accuracy: 32.68%
Epoch 4, Loss: 2.116175899725131, Accuracy: 34.01%
Test Accuracy: 36.88%
Epoch 5, Loss: 2.0800950630851416, Accuracy: 37.986%
Test Accuracy: 39.23%
Epoch 6, Loss: 2.0543632992088336, Accuracy: 40.438%
Test Accuracy: 39.71%
Epoch 7, Loss: 2.033445654317851, Accuracy: 42.658%
Test Accuracy: 42.61%
Epoch 8, Loss: 2.0150980592688637, Accuracy: 44.502%
Test Accuracy: 44.79%
Epoch 9, Loss: 1.9956189235457984, Accuracy: 46.558%
Test Accuracy: 43.75%
Epoch 10, Loss: 1.9786441292604218, Accuracy: 48.198%
Test Accuracy: 46.88%



```

## ▼ Plotting

```
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.title("Training Accuracy")
plt.plot(train_accuracies)
plt.xlabel("Epoch")
plt.ylabel("Accuracy (%)")
```

```
plt.subplot(1, 2, 2)
plt.title("Training Loss")
plt.plot(train_losses)
plt.xlabel("Epoch")
plt.ylabel("Loss")
```

```
plt.show()
```



### ▼ Extract Weights of the First Convolutional Layer (FCL)

```
weights = model.conv1.weight.data.numpy()
```

### ▼ Visualize FCL's Weights (assuming 32 filter maps)




```

for i in range(32):
    plt.subplot(4, 8, i+1)
    plt.imshow(weights[i][0], cmap='gray')
    plt.axis('off')

plt.tight_layout()
plt.show()

```



## ▼ Problem 2

The paper "Visualizing and Understanding Convolutional Networks" by Matthew Zeiler and Rob Fergus is a seminal work in deep learning. Published in 2013, it has significantly demystified the often "black-box" nature of Convolutional Neural Networks (CNNs), offering academics and industry practitioners critical insights into the architecture and function of CNNs.

### ▼ Key Ideas:

1. **Deconvolutional Network:** A pivotal contribution of this paper is introducing a "deconvolutional network," a specialized architecture designed to map feature activations back to the input pixel space. This deconvolutional network effectively reverses the operations of the CNN's convolutional layers. The deconvolutional network is an invaluable tool for visualizing and understanding what each layer in the CNN has learned, thereby elucidating which features are most salient for the network's performance.
2. **Layer-wise Visualization:** The deconvolutional network enables the authors to showcase the hierarchical nature of the features learned by a CNN. Lower layers tend to capture primary attributes such as edges and textures. In contrast, higher layers capture increasingly abstract and complex combinations of the features learned in the preceding layers, which could range from object parts to entire objects.
3. **Feature Abstractions:** The paper demonstrates that feature abstractions become progressively complex as one moves deeper into the network. For example, while the first layer may specialize in edge detection, subsequent layers might focus on shapes by aggregating these edges, and even deeper layers may identify intricate structures or objects.
4. **Sensitivity Analysis:** This technique is introduced to pinpoint the input stimuli that activate specific feature maps, providing an alternative avenue for interpreting what the network is learning. Sensitivity analysis is crucial not just for understanding but also for diagnosing the model's behavior, highlighting how minute changes in the input can have a significant impact on the output.
5. **Improving Performance:** The visualization techniques are leveraged to identify and rectify architectural and training regimen issues, resulting in an optimized CNN model. This enhanced model sets new performance benchmarks across several standard metrics and challenges, including the ImageNet competition.
6. **Fine-Tuning and Transfer Learning:** The authors empirically demonstrate that the features learned in one domain or task can be effectively transferred to different but related fields or jobs. This domain transfer underscores the robust generalizability of CNNs.
7. **Validation:** The rigorous validation process credibly establishes the efficacy of the introduced visualization techniques. Performance comparisons before and after implementing architectural insights derived from visualizations lend additional credibility to these techniques.

The paper provides a comprehensive set of tools and methodologies for understanding, interpreting, and even diagnosing CNNs. These contributions have led to performance improvements and built greater trust in deploying CNNs across various applications. Moreover, the techniques introduced have set a precedent for future work in CNN interpretability and have been widely adopted across academic and industrial settings. Since its publication, the paper has been a cornerstone in deep learning, inspiring subsequent research on making neural networks more interpretable and transparent.

▼ Applying one technique discussed in paper (Single Layer Visualization)

```
def plot_kernels(tensor, num_cols=8):
    # make 4D tensor if it is 3D
    if len(tensor.shape) == 3:
        tensor = tensor.unsqueeze(1)
    num_kernels = tensor.shape[0]
    num_rows = 1 + num_kernels // num_cols
    fig = plt.figure(figsize=(num_cols, num_rows))
    for i in range(num_kernels):
        ax1 = fig.add_subplot(num_rows, num_cols, i + 1)
        ax1.imshow(tensor[i][0, :, :], cmap="gray")
        ax1.axis('off')
        ax1.set_xticklabels([])
        ax1.set_yticklabels([])

    plt.subplots_adjust(wspace=0.1, hspace=0.1)
    plt.show()
```

```
def visualize_layer(model, layer_name, input_image):
    activation = {}
    def hook_fn(module, input, output):
        activation[layer_name] = output.detach()

    # Register hook
    layer = dict(*model.named_modules())[layer_name]
    hook = layer.register_forward_hook(hook_fn)

    # Forward pass (assuming the input image is already a 4D tensor)
    model(input_image.unsqueeze(0))


    # Remove hook
    hook.remove()

    act = activation[layer_name].squeeze()
    # Make the activation 3D if it is 2D
    if len(act.shape) == 2:
        act = act.unsqueeze(0)
    plot_kernels(act)
```

```
# Assuming `images` and `labels` are one batch from your test_loader
with torch.no_grad():
    for i, data in enumerate(test_loader, 0):
        images, labels = data
        # Just take the first image from the first batch
        img = images[0]

        # Visualize the activations of the first convolution layer
        visualize_layer(model, 'conv1', img)

    # Exit after first batch
    break
```



```

# Register hooks for Conv layers
for name, layer in model.named_modules():
    if isinstance(layer, nn.Conv2d):
        layer.register_forward_hook(hook_fn)

# In your test loop, after you get the model output
with torch.no_grad():
    for i, data in enumerate(test_loader, 0):
        images, labels = data
        outputs = model(images) # This will also populate 'activations' due to hook

        # Log activations from convolutional layers
        grid = torchvision.utils.make_grid(images)
        writer.add_image('images', grid, 0)
        writer.add_graph(model, images)
        for name, activation in activations.items():
            writer.add_histogram(f"{name}.activation", activation.flatten(), i)

<ipython-input-7-52bdba1e040d>:12: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['conv1_weights'] = self.conv1.weight.data.cpu().numpy()
<ipython-input-7-52bdba1e040d>:13: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['conv1_biases'] = self.conv1.bias.data.cpu().numpy()
<ipython-input-7-52bdba1e040d>:16: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['conv2_weights'] = self.conv2.weight.data.cpu().numpy()
<ipython-input-7-52bdba1e040d>:17: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['conv2_biases'] = self.conv2.bias.data.cpu().numpy()
<ipython-input-7-52bdba1e040d>:21: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['fc1_weights'] = self.fc1.weight.data.cpu().numpy()
<ipython-input-7-52bdba1e040d>:22: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['fc1_biases'] = self.fc1.bias.data.cpu().numpy()
<ipython-input-7-52bdba1e040d>:24: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['fc2_weights'] = self.fc2.weight.data.cpu().numpy()
<ipython-input-7-52bdba1e040d>:25: TracerWarning: Converting a tensor to a numpy array. This operation is deprecated and will be removed in a future PyTorch version. Please use tensor.detach().cpu().numpy() instead.
self.stats['fc2_biases'] = self.fc2.bias.data.cpu().numpy()

```

## ▼ Problem 3

### ▼ Importing Libraries

```
from datetime import datetime
from pathlib import Path

import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

from torchvision import datasets, transforms
from torch.utils.data import DataLoader
from torch.utils.tensorboard import SummaryWriter
```

## ▼ Defining Models

```
class RNNBaseModel(nn.Module):
    def __init__(self, rnn_type, input_dim, hidden_dim, output_dim):
        super(RNNBaseModel, self).__init__()
        self.rnn = rnn_type(input_dim, hidden_dim, batch_first=True)
        self.fc = nn.Linear(hidden_dim, output_dim)
        self.stats = {}

    def forward(self, x):
        output, hn = self.rnn(x)
        self.stats['rnn_weights'] = self.rnn.weight_hh_l0.data.cpu().numpy()
        self.stats['rnn_biases'] = self.rnn.bias_hh_l0.data.cpu().numpy()

        output = self.fc(output[-1, :, :])
        self.stats['fc_weights'] = self.fc.weight.data.cpu().numpy()
        self.stats['fc_biases'] = self.fc.bias.data.cpu().numpy()

        output = F.log_softmax(output, dim=1)
        return output
```

```

class AdaptiveCNNModel(nn.Module):
    def __init__(self, num_filters1, num_filters2, fc_size):
        super(AdaptiveCNNModel, self).__init__()
        self.conv1 = nn.Conv2d(1, num_filters1, 5)
        self.conv2 = nn.Conv2d(num_filters1, num_filters2, 5)
        self.fc1_size = num_filters2 * 4 * 4 # This needs to be calculated based on
        self.fc1 = nn.Linear(self.fc1_size, fc_size)
        self.fc2 = nn.Linear(fc_size, 10)
        self.stats = {}

    def forward(self, x):
        # First Conv Layer
        x = self.conv1(x)
        self.stats['conv1_weights'] = self.conv1.weight.data.cpu().numpy()
        self.stats['conv1_biases'] = self.conv1.bias.data.cpu().numpy()
        x = F.relu(x)
        x = F.max_pool2d(x, 2)

        # Second Conv Layer
        x = self.conv2(x)
        self.stats['conv2_weights'] = self.conv2.weight.data.cpu().numpy()
        self.stats['conv2_biases'] = self.conv2.bias.data.cpu().numpy()
        x = F.relu(x)
        x = F.max_pool2d(x, 2)

        # Fully Connected Layers
        x = x.view(x.size(0), -1) # Flattening
        x = self.fc1(x)
        self.stats['fc1_weights'] = self.fc1.weight.data.cpu().numpy()
        self.stats['fc1_biases'] = self.fc1.bias.data.cpu().numpy()
        x = F.relu(x)

        x = self.fc2(x)
        self.stats['fc2_weights'] = self.fc2.weight.data.cpu().numpy()
        self.stats['fc2_biases'] = self.fc2.bias.data.cpu().numpy()

        x = F.log_softmax(x, dim=1)

        return x

```

## ▼ Declaring Parameters



## ▼ General Hyperparams

```
batch_size = 28
logging_interval = 10
learning_rates = [0.001, 0.01]
optimizers = [optim.Adam, optim.SGD]
```

## ▼ RNN Hyperparams

```
hidden_dims = [64, 128]
```

## ▼ CNN Hyperparams

```
fc_size_list = [128, 256]
num_filters1_list = [32, 64]
num_filters2_list = [64, 128]
```

## ▼ Data Loading

```
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5, ),

trainset = datasets.MNIST(root='./data', train=True, download=True, transform=trans
testset = datasets.MNIST(root='./data', train=False, download=True, transform=trans

train_loader = DataLoader(trainset, batch_size=batch_size, shuffle=True, drop_last=
test_loader = DataLoader(testset, batch_size=batch_size, shuffle=False, drop_last=1
```

## ▼ Hyperparameter Tuning & Visuals

### ▼ RNN

```
def timing_decorator(func):
    def timing_wrapper(*args, **kwargs):
        start = time.time()
        result = func(*args, **kwargs)
        end = time.time()
        print(f"Timing for {func.__name__}: {end - start} seconds")
        return result
    return timing_wrapper
```

```
datetime_str = datetime.now().strftime('%d-%m-%Y_%H-%M-%S')
```

```
base_folder = Path('runs/rnn_experiment')  
base_folder.mkdir(parents=True, exist_ok=True)
```

```
logging_dir = base_folder / Path(datetime_str)  
logging_dir.mkdir(exist_ok=True)
```

```
logging_dir = str(logging_dir.absolute())
```

```
writer = SummaryWriter(log_dir=logging_dir)
```

```
%tensorboard --logdir {logging_dir} --port 7001
```

Filter runs (regex)

Filter tags (regex)

- All
- Scalars
- Image
- Histogram



Run




.




Loss

Loss/train\_batch



histogram 10 cards

histogram/fc.bias



```
train_losses = []
test accuracies = []
train accuracies = []
```

```

for i, hidden_dim in enumerate(hidden_dims):
    for j, lr in enumerate(learning_rates):
        for k, opt in enumerate(optimizers):
            model = RNNBaseModel(nn.RNN, batch_size, hidden_dim, 10)
            criterion = nn.CrossEntropyLoss()
            optimizer = opt(model.parameters(), lr=lr)
            print(f'Hyperparameters: [{i+1} hidden_dim={hidden_dim}, {j+1} lr={lr},
            for epoch in range(2): # Limiting to 2 epochs for demonstration
                running_loss = 0.0
                correct = 0
                model.train()
                for batch_idx, (images, labels) in enumerate(train_loader):
                    n_iter = epoch * len(train_loader) + batch_idx
                    images = images.view(-1, batch_size, batch_size)
                    outputs = model(images)
                    loss = criterion(outputs, labels)

                    optimizer.zero_grad()
                    loss.backward()
                    optimizer.step()

                    total += labels.size(0)
                    running_loss += loss.item()
                    _, predicted = torch.max(outputs.data, 1)
                    correct += (predicted == labels).sum().item()

                if (batch_idx + 1) % logging_interval == 0:
                    writer.add_scalar('Loss/train_batch', loss.item(), n_iter)
                    writer.add_scalar('Accuracy/train', 100. * correct / total, ep

            for key, value in model.stats.items():
                value_flat = value.flatten()
                writer.add_scalar(f'statistics/{key}_mean', np.mean(value_flat))
                writer.add_scalar(f'statistics/{key}_std', np.std(value_flat))
                writer.add_scalar(f'statistics/{key}_min', np.min(value_flat))
                writer.add_scalar(f'statistics/{key}_max', np.max(value_flat))
                writer.add_histogram(f'histogram/{key}', value_flat, n_iter)

            for name, param in model.named_parameters():
                if 'weight' in name:
                    writer.add_scalar(f'statistics/{name}_min', param.min())
                    writer.add_scalar(f'statistics/{name}_max', param.max())
                    writer.add_scalar(f'statistics/{name}_mean', param.mean())
                    writer.add_scalar(f'statistics/{name}_std', param.std())
                    writer.add_histogram(f'histogram/{name}', param, n_iter)
                elif 'bias' in name:

```

```

        writer.add_scalar(f'statistics/{name}_min', param.min)
        writer.add_scalar(f'statistics/{name}_max', param.max)
        writer.add_scalar(f'statistics/{name}_mean', param.mean)
        writer.add_scalar(f'statistics/{name}_std', param.std)
        writer.add_histogram(f'histogram/{name}', param, n_iter)

    print(f'Epoch [{epoch+1}/2], Step [{batch_idx+1}/{len(train_loader)}]')
    train_accuracies.append(100. * correct / total)
    train_losses.append(running_loss / total)

# Test loop
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for j, data in enumerate(test_loader):
        images, labels = data
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
    m_iter = epoch * len(test_loader) + j
    if (j + 1) % logging_interval == 0:
        for key, value in model.stats.items():
            value_flat = value.flatten()
            writer.add_scalar(f'statistics/test_{key}_mean', np.mean(value_flat))
            writer.add_scalar(f'statistics/test_{key}_std', np.std(value_flat))
            writer.add_scalar(f'statistics/test_{key}_min', np.min(value_flat))
            writer.add_scalar(f'statistics/test_{key}_max', np.max(value_flat))
            writer.add_histogram(f'histogram/test_{key}', value_flat)

    for name, param in model.named_parameters():
        if 'weight' in name:
            writer.add_scalar(f'statistics/test_{name}_min', param.min)
            writer.add_scalar(f'statistics/test_{name}_max', param.max)
            writer.add_scalar(f'statistics/test_{name}_mean', param.mean)
            writer.add_scalar(f'statistics/test_{name}_std', param.std)
            writer.add_histogram(f'histogram/test_{name}', param, n_iter)
        elif 'bias' in name:
            writer.add_scalar(f'statistics/test_{name}_min', param.min)
            writer.add_scalar(f'statistics/test_{name}_max', param.max)
            writer.add_scalar(f'statistics/test_{name}_mean', param.mean)
            writer.add_scalar(f'statistics/test_{name}_std', param.std)
            writer.add_histogram(f'histogram/test_{name}', param, n_iter)
    writer.add_scalar('Accuracy/test', 100. * correct / total, epoch)
    print(f"Test Accuracy: {100 * correct / total}%")
    test_accuracies.append(100. * correct / total)

```

writer.close()

Epoch	[2/2]	, Step	[1360/2142]	, Loss:	2.2967
Epoch	[2/2]	, Step	[1370/2142]	, Loss:	2.3038
Epoch	[2/2]	, Step	[1380/2142]	, Loss:	2.3018
Epoch	[2/2]	, Step	[1390/2142]	, Loss:	2.3152
Epoch	[2/2]	, Step	[1400/2142]	, Loss:	2.2866
Epoch	[2/2]	, Step	[1410/2142]	, Loss:	2.3069
Epoch	[2/2]	, Step	[1420/2142]	, Loss:	2.3122
Epoch	[2/2]	, Step	[1430/2142]	, Loss:	2.2868
Epoch	[2/2]	, Step	[1440/2142]	, Loss:	2.3032
Epoch	[2/2]	, Step	[1450/2142]	, Loss:	2.3119
Epoch	[2/2]	, Step	[1460/2142]	, Loss:	2.3070
Epoch	[2/2]	, Step	[1470/2142]	, Loss:	2.3005
Epoch	[2/2]	, Step	[1480/2142]	, Loss:	2.2944
Epoch	[2/2]	, Step	[1490/2142]	, Loss:	2.2893
Epoch	[2/2]	, Step	[1500/2142]	, Loss:	2.2937
Epoch	[2/2]	, Step	[1510/2142]	, Loss:	2.3032
Epoch	[2/2]	, Step	[1520/2142]	, Loss:	2.3190
Epoch	[2/2]	, Step	[1530/2142]	, Loss:	2.3030
Epoch	[2/2]	, Step	[1540/2142]	, Loss:	2.3203
Epoch	[2/2]	, Step	[1550/2142]	, Loss:	2.2840
Epoch	[2/2]	, Step	[1560/2142]	, Loss:	2.2923
Epoch	[2/2]	, Step	[1570/2142]	, Loss:	2.3305
Epoch	[2/2]	, Step	[1580/2142]	, Loss:	2.3138
Epoch	[2/2]	, Step	[1590/2142]	, Loss:	2.3061
Epoch	[2/2]	, Step	[1600/2142]	, Loss:	2.2852
Epoch	[2/2]	, Step	[1610/2142]	, Loss:	2.2641
Epoch	[2/2]	, Step	[1620/2142]	, Loss:	2.3142
Epoch	[2/2]	, Step	[1630/2142]	, Loss:	2.2976
Epoch	[2/2]	, Step	[1640/2142]	, Loss:	2.3201
Epoch	[2/2]	, Step	[1650/2142]	, Loss:	2.2943
Epoch	[2/2]	, Step	[1660/2142]	, Loss:	2.3003
Epoch	[2/2]	, Step	[1670/2142]	, Loss:	2.3002
Epoch	[2/2]	, Step	[1680/2142]	, Loss:	2.3235
Epoch	[2/2]	, Step	[1690/2142]	, Loss:	2.3081
Epoch	[2/2]	, Step	[1700/2142]	, Loss:	2.3208
Epoch	[2/2]	, Step	[1710/2142]	, Loss:	2.3099
Epoch	[2/2]	, Step	[1720/2142]	, Loss:	2.2969
Epoch	[2/2]	, Step	[1730/2142]	, Loss:	2.3031
Epoch	[2/2]	, Step	[1740/2142]	, Loss:	2.2741
Epoch	[2/2]	, Step	[1750/2142]	, Loss:	2.3224
Epoch	[2/2]	, Step	[1760/2142]	, Loss:	2.2882
Epoch	[2/2]	, Step	[1770/2142]	, Loss:	2.2917
Epoch	[2/2]	, Step	[1780/2142]	, Loss:	2.2969
Epoch	[2/2]	, Step	[1790/2142]	, Loss:	2.2984
Epoch	[2/2]	, Step	[1800/2142]	, Loss:	2.3042
Epoch	[2/2]	, Step	[1810/2142]	, Loss:	2.3175
Epoch	[2/2]	, Step	[1820/2142]	, Loss:	2.2907
Epoch	[2/2]	, Step	[1830/2142]	, Loss:	2.2841
Epoch	[2/2]	, Step	[1840/2142]	, Loss:	2.3157
Epoch	[2/2]	, Step	[1850/2142]	, Loss:	2.3179
Epoch	[2/2]	, Step	[1860/2142]	, Loss:	2.2995

```
Epoch [2/2], Step [1870/2142], Loss: 2.3148
Epoch [2/2], Step [1880/2142], Loss: 2.2994
Epoch [2/2], Step [1890/2142], Loss: 2.2845
Epoch [2/2], Step [1900/2142], Loss: 2.2932
Epoch [2/2], Step [1910/2142], Loss: 2.3004
Epoch [2/2], Step [1920/2142], Loss: 2.3109
Epoch [2/2], Step [1930/2142], Loss: 2.3233
Epoch [2/2], Step [1940/2142], Loss: 2.3079
```

## ▼ LSTM

```
datetime_str = datetime.now().strftime('%b%d_%H-%M-%S')
```

```
base_folder = Path('runs/lstm_experiment')
base_folder.mkdir(parents=True, exist_ok=True)
```

```
logging_dir = base_folder / Path(datetime_str)
logging_dir.mkdir(exist_ok=True)
```

```
logging_dir = str(logging_dir.absolute())
```

```
writer = SummaryWriter(log_dir=logging_dir)
```

```
%tensorboard --logdir {logging_dir} --port 7003
```

```
train_losses = []
test_accuracies = []
train_accuracies = []
```



```

for i, hidden_dim in enumerate(hidden_dims):
    for j, lr in enumerate(learning_rates):
        for k, opt in enumerate(optimizers):
            model = RNNBaseModel(nn.LSTM, batch_size, hidden_dim, 10)
            criterion = nn.CrossEntropyLoss()
            optimizer = opt(model.parameters(), lr=lr)
            print(f'Hyperparameters: [{i+1} hidden_dim={hidden_dim}, #{j+1} lr={lr}')
            for epoch in range(2): # Limiting to 2 epochs for demonstration
                running_loss = 0.0
                correct = 0
                model.train()
                for batch_idx, (images, labels) in enumerate(train_loader):
                    n_iter = epoch * len(train_loader) + batch_idx
                    images = images.view(-1, batch_size, batch_size)
                    outputs = model(images)
                    loss = criterion(outputs, labels)

                    optimizer.zero_grad()
                    loss.backward()
                    optimizer.step()

                    total += labels.size(0)
                    running_loss += loss.item()
                    _, predicted = torch.max(outputs.data, 1)
                    correct += (predicted == labels).sum().item()

                if (batch_idx + 1) % logging_interval == 0:
                    writer.add_scalar('Loss/train_batch', loss.item(), n_iter)
                    writer.add_scalar('Accuracy/train', 100. * correct / total, n_iter)

                for key, value in model.stats.items():
                    value_flat = value.flatten()
                    writer.add_scalar(f'statistics/train_{key}_mean', np.mean(value_flat))
                    writer.add_scalar(f'statistics/train_{key}_std', np.std(value_flat))
                    writer.add_scalar(f'statistics/train_{key}_min', np.min(value_flat))
                    writer.add_scalar(f'statistics/train_{key}_max', np.max(value_flat))
                    writer.add_histogram(f'histogram/train_{key}', value_flat)

            for name, param in model.named_parameters():
                if 'weight' in name:
                    writer.add_scalar(f'statistics/train_{name}_min', param.min())
                    writer.add_scalar(f'statistics/train_{name}_max', param.max())
                    writer.add_scalar(f'statistics/train_{name}_mean', param.mean())
                    writer.add_scalar(f'statistics/train_{name}_std', param.std())
                    writer.add_histogram(f'histogram/train_{name}', param.data)
                elif 'bias' in name:
                    writer.add_scalar(f'statistics/train_{name}_min', param.min())

```

```

        writer.add_scalar(f'statistics/train_{name}_max', par
        writer.add_scalar(f'statistics/train_{name}_mean', pa
        writer.add_scalar(f'statistics/train_{name}_std', par
        writer.add_histogram(f'histogram/train_{name}', paran

    print(f'Epoch [{epoch+1}/2], Step [{batch_idx+1}/{len(train_l
train_accuracies.append(100. * correct / total)
train_losses.append(running_loss / total)

# Test Loop
model.eval()
correct = 0
total = 0

with torch.no_grad():
    for j, data in enumerate(test_loader):

        images, labels = data
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        m_iter = epoch * len(test_loader) + j
        if (j + 1) % logging_interval == 0:
            for key, value in model.stats.items():
                value_flat = value.flatten()
                writer.add_scalar(f'statistics/test_{key}_mean', np
                writer.add_scalar(f'statistics/test_{key}_std', np.
                writer.add_scalar(f'statistics/test_{key}_min', np.
                writer.add_scalar(f'statistics/test_{key}_max', np.
                writer.add_histogram(f'histogram/test_{key}', value

    for name, param in model.named_parameters():
        if 'weight' in name:
            writer.add_scalar(f'statistics/test_{name}_min'
            writer.add_scalar(f'statistics/test_{name}_max'
            writer.add_scalar(f'statistics/test_{name}_mean'
            writer.add_scalar(f'statistics/test_{name}_std'
            writer.add_histogram(f'histogram/test_{name}',
        elif 'bias' in name:
            writer.add_scalar(f'statistics/test_{name}_min'
            writer.add_scalar(f'statistics/test_{name}_max'
            writer.add_scalar(f'statistics/test_{name}_mean'
            writer.add_scalar(f'statistics/test_{name}_std'
            writer.add_histogram(f'histogram/test_{name}',
    writer.add_scalar('Accuracy/test', 100 * correct / total, epoch)
    print(f"Test Accuracy: {100 * correct / total}%")
    test_accuracies.append(100 * correct / total)

```

```
writer.close()
```

```
Epoch [2/2], Step [1000/2142], Loss: 2.3005
Epoch [2/2], Step [1010/2142], Loss: 2.3205
Epoch [2/2], Step [1020/2142], Loss: 2.2924
Epoch [2/2], Step [1030/2142], Loss: 2.2869
Epoch [2/2], Step [1040/2142], Loss: 2.2891
Epoch [2/2], Step [1050/2142], Loss: 2.3053
Epoch [2/2], Step [1060/2142], Loss: 2.2917
Epoch [2/2], Step [1070/2142], Loss: 2.2815
Epoch [2/2], Step [1080/2142], Loss: 2.2734
Epoch [2/2], Step [1090/2142], Loss: 2.2834
Epoch [2/2], Step [1100/2142], Loss: 2.3043
Epoch [2/2], Step [1110/2142], Loss: 2.3071
Epoch [2/2], Step [1120/2142], Loss: 2.3112
Epoch [2/2], Step [1130/2142], Loss: 2.3069
Epoch [2/2], Step [1140/2142], Loss: 2.3059
Epoch [2/2], Step [1150/2142], Loss: 2.2956
Epoch [2/2], Step [1160/2142], Loss: 2.2901
Epoch [2/2], Step [1170/2142], Loss: 2.3133
Epoch [2/2], Step [1180/2142], Loss: 2.3049
Epoch [2/2], Step [1190/2142], Loss: 2.3001
Epoch [2/2], Step [1200/2142], Loss: 2.3021
Epoch [2/2], Step [1210/2142], Loss: 2.2963
Epoch [2/2], Step [1220/2142], Loss: 2.2869
Epoch [2/2], Step [1230/2142], Loss: 2.3057
Epoch [2/2], Step [1240/2142], Loss: 2.3058
Epoch [2/2], Step [1250/2142], Loss: 2.3005
Epoch [2/2], Step [1260/2142], Loss: 2.2998
Epoch [2/2], Step [1270/2142], Loss: 2.3089
Epoch [2/2], Step [1280/2142], Loss: 2.2985
Epoch [2/2], Step [1290/2142], Loss: 2.2975
Epoch [2/2], Step [1300/2142], Loss: 2.2971
Epoch [2/2], Step [1310/2142], Loss: 2.3132
Epoch [2/2], Step [1320/2142], Loss: 2.2831
Epoch [2/2], Step [1330/2142], Loss: 2.3156
Epoch [2/2], Step [1340/2142], Loss: 2.3099
Epoch [2/2], Step [1350/2142], Loss: 2.2891
Epoch [2/2], Step [1360/2142], Loss: 2.3203
Epoch [2/2], Step [1370/2142], Loss: 2.2906
Epoch [2/2], Step [1380/2142], Loss: 2.3050
Epoch [2/2], Step [1390/2142], Loss: 2.3028
Epoch [2/2], Step [1400/2142], Loss: 2.2939
Epoch [2/2], Step [1410/2142], Loss: 2.3006
Epoch [2/2], Step [1420/2142], Loss: 2.2987
Epoch [2/2], Step [1430/2142], Loss: 2.3054
Epoch [2/2], Step [1440/2142], Loss: 2.3161
Epoch [2/2], Step [1450/2142], Loss: 2.3000
Epoch [2/2], Step [1460/2142], Loss: 2.3095
Epoch [2/2], Step [1470/2142], Loss: 2.2944
Epoch [2/2], Step [1480/2142], Loss: 2.3143
Epoch [2/2], Step [1490/2142], Loss: 2.2916
```

```
Epoch [2/2], Step [1500/2142], Loss: 2.2925
Epoch [2/2], Step [1510/2142], Loss: 2.3234
Epoch [2/2], Step [1520/2142], Loss: 2.2966
Epoch [2/2], Step [1530/2142], Loss: 2.2837
Epoch [2/2], Step [1540/2142], Loss: 2.2856
Epoch [2/2], Step [1550/2142], Loss: 2.2823
Epoch [2/2], Step [1560/2142], Loss: 2.3100
Epoch [2/2], Step [1570/2142], Loss: 2.2963
Epoch [2/2], Step [1580/2142], Loss: 2.2991
```

## ▼ GRU

```
datetime_str = datetime.now().strftime('%b%d_%H-%M-%S')
```

```
base_folder = Path('runs/gru_experiment')
base_folder.mkdir(parents=True, exist_ok=True)
```

```
logging_dir = base_folder / Path(datetime_str)
logging_dir.mkdir(exist_ok=True)
```

```
logging_dir = str(logging_dir.absolute())
```

```
writer = SummaryWriter(log_dir=logging_dir)
```

```
%tensorboard --logdir {logging_dir} --port 7005
```

```
train_losses = []
test_accuracies = []
train_accuracies = []
```

```

for i, hidden_dim in enumerate(hidden_dims):
    for j, lr in enumerate(learning_rates):
        for k, opt in enumerate(optimizers):
            model = RNNBaseModel(nn.GRU, batch_size, hidden_dim, 10)
            criterion = nn.CrossEntropyLoss()
            optimizer = opt(model.parameters(), lr=lr)
            print(f'Hyperparameters: [{i+1} hidden_dim={hidden_dim}, #{j+1} lr={lr}')
            for epoch in range(2): # Limiting to 2 epochs for demonstration
                running_loss = 0.0
                correct = 0
                model.train()
                for batch_idx, (images, labels) in enumerate(train_loader):
                    n_iter = epoch * len(train_loader) + batch_idx
                    images = images.view(-1, batch_size, batch_size)
                    outputs = model(images)
                    loss = criterion(outputs, labels)

                    optimizer.zero_grad()
                    loss.backward()
                    optimizer.step()

                    total += labels.size(0)
                    running_loss += loss.item()
                    _, predicted = torch.max(outputs.data, 1)
                    correct += (predicted == labels).sum().item()

                if (batch_idx + 1) % logging_interval == 0:
                    writer.add_scalar('Loss/train_batch', loss.item(), n_iter)
                    writer.add_scalar('Accuracy/train', 100 * correct / total,

                    for key, value in model.stats.items():
                        value_flat = value.flatten()
                        writer.add_scalar(f'statistics/train_{key}_mean', np.mean(value_flat))
                        writer.add_scalar(f'statistics/train_{key}_std', np.std(value_flat))
                        writer.add_scalar(f'statistics/train_{key}_min', np.min(value_flat))
                        writer.add_scalar(f'statistics/train_{key}_max', np.max(value_flat))
                        writer.add_histogram(f'histogram/train_{key}', value_flat)

                    for name, param in model.named_parameters():
                        if 'weight' in name:
                            writer.add_scalar(f'statistics/train_{name}_min', np.min(param.data))
                            writer.add_scalar(f'statistics/train_{name}_max', np.max(param.data))
                            writer.add_scalar(f'statistics/train_{name}_mean', np.mean(param.data))
                            writer.add_scalar(f'statistics/train_{name}_std', np.std(param.data))
                            writer.add_histogram(f'histogram/train_{name}', param.data)
                        elif 'bias' in name:
                            writer.add_scalar(f'statistics/train_{name}_min', np.min(param.data))

```

```

        writer.add_scalar(f'statistics/train_{name}_max', p
        writer.add_scalar(f'statistics/train_{name}_mean',
        writer.add_scalar(f'statistics/train_{name}_std', p
        writer.add_histogram(f'histogram/train_{name}', par

    print(f'Epoch [{epoch+1}/2], Step [{batch_idx+1}/{len(trair
train_accuracies.append(100. * correct / total)
train_losses.append(running_loss / total)

# Test Loop
model.eval()
correct = 0
total = 0

with torch.no_grad():
    for j, data in enumerate(test_loader):

        images, labels = data
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        m_iter = epoch * len(test_loader) + j
        if (j + 1) % logging_interval == 0:
            for key, value in model.stats.items():
                value_flat = value.flatten()
                writer.add_scalar(f'statistics/test_{key}_mean',
                writer.add_scalar(f'statistics/test_{key}_std', r
                writer.add_scalar(f'statistics/test_{key}_min', r
                writer.add_scalar(f'statistics/test_{key}_max', r
                writer.add_histogram(f'histogram/test_{key}', val

        for name, param in model.named_parameters():
            if 'weight' in name:
                writer.add_scalar(f'statistics/test_{name}_mi
                writer.add_scalar(f'statistics/test_{name}_ma
                writer.add_scalar(f'statistics/test_{name}_me
                writer.add_scalar(f'statistics/test_{name}_st
                writer.add_histogram(f'histogram/test_{name}')
            elif 'bias' in name:
                writer.add_scalar(f'statistics/test_{name}_mi
                writer.add_scalar(f'statistics/test_{name}_ma
                writer.add_scalar(f'statistics/test_{name}_me
                writer.add_scalar(f'statistics/test_{name}_st
                writer.add_histogram(f'histogram/test_{name}')
        writer.add_scalar('Accuracy/test', 100 * correct / total, epoch
        print(f"Test Accuracy: {100 * correct / total}%")
        test_accuracies.append(100 * correct / total)

```

```
writer.close()
```

```
Epoch [2/2], Step [1270/2142], Loss: 2.2935
Epoch [2/2], Step [1280/2142], Loss: 2.3112
Epoch [2/2], Step [1290/2142], Loss: 2.2980
Epoch [2/2], Step [1300/2142], Loss: 2.2978
Epoch [2/2], Step [1310/2142], Loss: 2.2916
Epoch [2/2], Step [1320/2142], Loss: 2.3099
Epoch [2/2], Step [1330/2142], Loss: 2.3106
Epoch [2/2], Step [1340/2142], Loss: 2.3031
Epoch [2/2], Step [1350/2142], Loss: 2.2972
Epoch [2/2], Step [1360/2142], Loss: 2.3043
Epoch [2/2], Step [1370/2142], Loss: 2.3043
Epoch [2/2], Step [1380/2142], Loss: 2.2928
Epoch [2/2], Step [1390/2142], Loss: 2.3021
Epoch [2/2], Step [1400/2142], Loss: 2.3011
Epoch [2/2], Step [1410/2142], Loss: 2.2874
Epoch [2/2], Step [1420/2142], Loss: 2.2883
Epoch [2/2], Step [1430/2142], Loss: 2.2764
Epoch [2/2], Step [1440/2142], Loss: 2.3087
Epoch [2/2], Step [1450/2142], Loss: 2.3034
Epoch [2/2], Step [1460/2142], Loss: 2.3235
Epoch [2/2], Step [1470/2142], Loss: 2.2640
Epoch [2/2], Step [1480/2142], Loss: 2.2672
Epoch [2/2], Step [1490/2142], Loss: 2.3012
Epoch [2/2], Step [1500/2142], Loss: 2.2908
Epoch [2/2], Step [1510/2142], Loss: 2.3004
Epoch [2/2], Step [1520/2142], Loss: 2.3027
Epoch [2/2], Step [1530/2142], Loss: 2.3018
Epoch [2/2], Step [1540/2142], Loss: 2.3055
Epoch [2/2], Step [1550/2142], Loss: 2.3197
Epoch [2/2], Step [1560/2142], Loss: 2.3000
Epoch [2/2], Step [1570/2142], Loss: 2.3136
Epoch [2/2], Step [1580/2142], Loss: 2.3181
Epoch [2/2], Step [1590/2142], Loss: 2.3005
Epoch [2/2], Step [1600/2142], Loss: 2.3016
Epoch [2/2], Step [1610/2142], Loss: 2.2847
Epoch [2/2], Step [1620/2142], Loss: 2.2909
Epoch [2/2], Step [1630/2142], Loss: 2.3201
Epoch [2/2], Step [1640/2142], Loss: 2.2936
Epoch [2/2], Step [1650/2142], Loss: 2.3291
Epoch [2/2], Step [1660/2142], Loss: 2.2985
Epoch [2/2], Step [1670/2142], Loss: 2.2786
Epoch [2/2], Step [1680/2142], Loss: 2.2781
Epoch [2/2], Step [1690/2142], Loss: 2.2950
Epoch [2/2], Step [1700/2142], Loss: 2.2900
Epoch [2/2], Step [1710/2142], Loss: 2.3083
Epoch [2/2], Step [1720/2142], Loss: 2.2891
Epoch [2/2], Step [1730/2142], Loss: 2.2922
Epoch [2/2], Step [1740/2142], Loss: 2.2854
Epoch [2/2], Step [1750/2142], Loss: 2.3284
Epoch [2/2], Step [1760/2142], Loss: 2.3117
```



```
Epoch [2/2], Step [1770/2142], Loss: 2.3148
Epoch [2/2], Step [1780/2142], Loss: 2.2871
Epoch [2/2], Step [1790/2142], Loss: 2.2938
Epoch [2/2], Step [1800/2142], Loss: 2.3186
Epoch [2/2], Step [1810/2142], Loss: 2.2807
Epoch [2/2], Step [1820/2142], Loss: 2.2953
Epoch [2/2], Step [1830/2142], Loss: 2.3033
Epoch [2/2], Step [1840/2142], Loss: 2.3407
Epoch [2/2], Step [1850/2142], Loss: 2.3090
```

## ▼ CNN Comparison

```
datetime_str = datetime.now().strftime('%b%d_%H-%M-%S')
```

```
base_folder = Path('runs/cnn_experiment')
base_folder.mkdir(parents=True, exist_ok=True)
```

```
logging_dir = base_folder / Path(datetime_str)
logging_dir.mkdir(exist_ok=True)
```

```
logging_dir = str(logging_dir.absolute())
```

```
writer = SummaryWriter(log_dir=logging_dir)
```

```
%tensorboard --logdir {logging_dir} --port 7007
```

```
for i, num_filters1 in enumerate(num_filters1_list):
    for j, num_filters2 in enumerate(num_filters2_list):
        for k, fc_size in enumerate(fc_size_list):
            for l, lr in enumerate(learning_rates):
```

```

for m, opt in enumerate(optimizers):
    model = AdaptiveCNNModel(num_filters1, num_filters2, fc_size)
    cnn_optimizer = opt(model.parameters(), lr=lr)
    cnn_criterion = nn.CrossEntropyLoss()
    print(f'Hyperparameters: [{i+1}] num_filters1={num_filters1}, #
    for epoch in range(2): # Limiting to 2 epochs for demonstratio
        model.train()
        correct = 0
        running_loss = 0
        for batch_idx, (data, target) in enumerate(train_loader):
            cnn_optimizer.zero_grad()
            output = model(data)
            loss = F.nll_loss(output, target)
            loss.backward()
            cnn_optimizer.step()
            running_loss += loss.item()

        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

        n_iter = epoch * len(train_loader) + batch_idx
        if (batch_idx + 1) % logging_interval == 0:
            writer.add_scalar('Loss/train_batch', loss.item(),
            writer.add_scalar('Accuracy/train', 100 * correct /

            for key, value in model.stats.items():
                value_flat = value.flatten()
                writer.add_scalar(f'statistics/train_{key}_mean
                writer.add_scalar(f'statistics/train_{key}_std'
                writer.add_scalar(f'statistics/train_{key}_min'
                writer.add_scalar(f'statistics/train_{key}_max'
                writer.add_histogram(f'histogram/train_{key}',

            for name, param in model.named_parameters():
                if 'weight' in name:
                    writer.add_scalar(f'statistics/train_{name}
                    writer.add_scalar(f'statistics/train_{name}
                    writer.add_scalar(f'statistics/train_{name}
                    writer.add_scalar(f'statistics/train_{name}
                    writer.add_histogram(f'histogram/train_{nam
                elif 'bias' in name:
                    writer.add_scalar(f'statistics/rain_{name}_
                    writer.add_scalar(f'statisticsrain_{name}_m
                    writer.add_scalar(f'statistics/rain_{name}_
                    writer.add_scalar(f'statistics/rain_{name}_
                    writer.add_histogram(f'histogram/rain_{name}

    print(f'Epoch [{epoch+1}/2], Step [{batch_idx+1}/{l

```

```

train_accuracies.append(100 * correct / total)
train_losses.append(running_loss / total)

# Test Loop
model.eval()
correct = 0
total = 0

with torch.no_grad():
    for j, data in enumerate(test_loader):

        images, labels = data
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        m_iter = epoch * len(test_loader) + j
        if (j + 1) % logging_interval == 0:
            for key, value in model.stats.items():
                value_flat = value.flatten()
                writer.add_scalar(f'statistics/test_{key}', value_flat, m_iter)
                writer.add_scalar(f'statistics/test_{key}', value_flat, m_iter)
                writer.add_scalar(f'statistics/test_{key}', value_flat, m_iter)
                writer.add_scalar(f'statistics/test_{key}', value_flat, m_iter)
                writer.add_histogram(f'histogram/test_{key}', value_flat, m_iter)

            for name, param in model.named_parameters():
                if 'weight' in name:
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_histogram(f'histogram/test_{name}', param.data, m_iter)
                elif 'bias' in name:
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_scalar(f'statistics/test_{name}', param.data, m_iter)
                    writer.add_histogram(f'histogram/test_{name}', param.data, m_iter)

            writer.add_scalar('Accuracy/test', 100 * correct / total, m_iter)
            print(f"Test Accuracy: {100 * correct / total}%")
            test_accuracies.append(100 * correct / total)

```

```

writer.close()

```

```

Epoch [2/2], Step [1230/2142], Loss: 0.1224
Epoch [2/2], Step [1240/2142], Loss: 0.0110
Epoch [2/2], Step [1250/2142], Loss: 0.0514
Epoch [2/2], Step [1260/2142], Loss: 0.0958
Epoch [2/2], Step [1270/2142], Loss: 0.0364

```

Epoch [2/2], Step [1280/2142], Loss: 0.0514  
Epoch [2/2], Step [1290/2142], Loss: 0.0323  
Epoch [2/2], Step [1300/2142], Loss: 0.3162  
Epoch [2/2], Step [1310/2142], Loss: 0.0842  
Epoch [2/2], Step [1320/2142], Loss: 0.0295  
Epoch [2/2], Step [1330/2142], Loss: 0.3445  
Epoch [2/2], Step [1340/2142], Loss: 0.0933  
Epoch [2/2], Step [1350/2142], Loss: 0.0151  
Epoch [2/2], Step [1360/2142], Loss: 0.0173  
Epoch [2/2], Step [1370/2142], Loss: 0.2274  
Epoch [2/2], Step [1380/2142], Loss: 0.3112  
Epoch [2/2], Step [1390/2142], Loss: 0.0226  
Epoch [2/2], Step [1400/2142], Loss: 0.0058  
Epoch [2/2], Step [1410/2142], Loss: 0.0115  
Epoch [2/2], Step [1420/2142], Loss: 0.0970  
Epoch [2/2], Step [1430/2142], Loss: 0.1925  
Epoch [2/2], Step [1440/2142], Loss: 0.0069  
Epoch [2/2], Step [1450/2142], Loss: 0.0170  
Epoch [2/2], Step [1460/2142], Loss: 0.0397  
Epoch [2/2], Step [1470/2142], Loss: 0.0341  
Epoch [2/2], Step [1480/2142], Loss: 0.1366  
Epoch [2/2], Step [1490/2142], Loss: 0.0322  
Epoch [2/2], Step [1500/2142], Loss: 0.2234  
Epoch [2/2], Step [1510/2142], Loss: 0.0547  
Epoch [2/2], Step [1520/2142], Loss: 0.0091  
Epoch [2/2], Step [1530/2142], Loss: 0.0106  
Epoch [2/2], Step [1540/2142], Loss: 0.0209  
Epoch [2/2], Step [1550/2142], Loss: 0.2992  
Epoch [2/2], Step [1560/2142], Loss: 0.3741  
Epoch [2/2], Step [1570/2142], Loss: 0.0175  
Epoch [2/2], Step [1580/2142], Loss: 0.0764  
Epoch [2/2], Step [1590/2142], Loss: 0.0872  
Epoch [2/2], Step [1600/2142], Loss: 0.0700  
Epoch [2/2], Step [1610/2142], Loss: 0.2254  
Epoch [2/2], Step [1620/2142], Loss: 0.0114  
Epoch [2/2], Step [1630/2142], Loss: 0.2529  
Epoch [2/2], Step [1640/2142], Loss: 0.0503  
Epoch [2/2], Step [1650/2142], Loss: 0.0189  
Epoch [2/2], Step [1660/2142], Loss: 0.0263  
Epoch [2/2], Step [1670/2142], Loss: 0.0956  
Epoch [2/2], Step [1680/2142], Loss: 0.0076  
Epoch [2/2], Step [1690/2142], Loss: 0.0230  
Epoch [2/2], Step [1700/2142], Loss: 0.0909  
Epoch [2/2], Step [1710/2142], Loss: 0.0690  
Epoch [2/2], Step [1720/2142], Loss: 0.0423  
Epoch [2/2], Step [1730/2142], Loss: 0.0151  
Epoch [2/2], Step [1740/2142], Loss: 0.0748  
Epoch [2/2], Step [1750/2142], Loss: 0.1872  
Epoch [2/2], Step [1760/2142], Loss: 0.0917  
Epoch [2/2], Step [1770/2142], Loss: 0.1895  
Epoch [2/2], Step [1780/2142], Loss: 0.0517  
Epoch [2/2], Step [1790/2142], Loss: 0.0204  
Epoch [2/2], Step [1800/2142], Loss: 0.0279

Epoch [2/2] Step [1910/2142] Loss: 0.0512